

What even is the `module` command?

Ed Bennett

@QuantumofEd



Swansea University
Prifysgol Abertawe

@SwanseaUni



SUPERCOMPUTING WALES
UWCHGYFRIFIADURA CYMRU

@SuperCompWales



@sa2c_swansea

SA2C Tech Chat, 2018-07-27

Slides: git.io/fNEcv

module **basics**

- HPC clusters provide a lot of software

module **basics**

- HPC clusters provide a lot of software (440 modules on SCW SWSB)

module **basics**

- HPC clusters provide a lot of software (440 modules on SCW SWSB)
- Not all of it can be in the environment at once

module **basics**

- HPC clusters provide a lot of software (440 modules on SCW SWSB)
- Not all of it can be in the environment at once
- Install each package in its own folder

module **basics**

- HPC clusters provide a lot of software (440 modules on SCW SWSB)
- Not all of it can be in the environment at once
- Install each package in its own folder
- Add to the \$PATH on request

Using module

- `module avail`

Using module

- `module avail`
- `module load`

Using module

- `module avail`
- `module load`
- `module list`

Using module

- `module avail`
- `module load`
- `module list`
- `module unload`

Using module

- `module avail`
- `module load`
- `module list`
- `module unload`
- `module purge`

Promises

Slides: git.io/fNEcv

Three thin, dark blue wavy lines that sweep across the bottom of the slide, starting from the left and ending on the right.

Promises

“Modules can be loaded and unloaded dynamically and atomically, in an clean fashion.”

Promises

“Modules can be loaded and unloaded dynamically and atomically, in an clean fashion.”

```
[s.e.j.bennett@c12 ~]$ echo $F90
```

Promises

“Modules can be loaded and unloaded dynamically and atomically, in an clean fashion.”

```
[s.e.j.bennett@c12 ~]$ echo $F90
```

```
[s.e.j.bennett@c12 ~]$ export F90=gfortran
```

Promises

“Modules can be loaded and unloaded dynamically and atomically, in an clean fashion.”

```
[s.e.j.bennett@c12 ~]$ echo $F90
```

```
[s.e.j.bennett@c12 ~]$ export F90=gfortran
```

```
[s.e.j.bennett@c12 ~]$ module load compiler/intel
```


Promises

“Modules can be loaded and unloaded dynamically and atomically, in an clean fashion.”

```
[s.e.j.bennett@c12 ~]$ echo $F90
```

```
[s.e.j.bennett@c12 ~]$ export F90=gfortran
```

```
[s.e.j.bennett@c12 ~]$ module load compiler/intel
```

```
[s.e.j.bennett@c12 ~]$ echo $F90
```

Promises

“Modules can be loaded and unloaded dynamically and atomically, in an clean fashion.”

```
[s.e.j.bennett@c12 ~]$ echo $F90
```

```
[s.e.j.bennett@c12 ~]$ export F90=gfortran
```

```
[s.e.j.bennett@c12 ~]$ module load compiler/intel
```

```
[s.e.j.bennett@c12 ~]$ echo $F90
```

```
ifort
```

```
[s.e.j.bennett@c12 ~]$ module unload compiler/intel
```

Promises

“Modules can be loaded and unloaded dynamically and atomically, in an clean fashion.”

```
[s.e.j.bennett@c12 ~]$ echo $F90
```

```
[s.e.j.bennett@c12 ~]$ export F90=gfortran
```

```
[s.e.j.bennett@c12 ~]$ module load compiler/intel
```

```
[s.e.j.bennett@c12 ~]$ echo $F90
```

```
ifort
```

```
[s.e.j.bennett@c12 ~]$ module unload compiler/intel
```

```
[s.e.j.bennett@c12 ~]$ echo $F90
```

Promises

“Modules can be loaded and unloaded dynamically and atomically, in an clean fashion.”

```
[s.e.j.bennett@c12 ~]$ echo $F90
```

```
[s.e.j.bennett@c12 ~]$ export F90=gfortran
```

```
[s.e.j.bennett@c12 ~]$ module load compiler/intel
```

```
[s.e.j.bennett@c12 ~]$ echo $F90
```

```
ifort
```

```
[s.e.j.bennett@c12 ~]$ module unload compiler/intel
```

```
[s.e.j.bennett@c12 ~]$ echo $F90
```

```
[s.e.j.bennett@c12 ~]$
```

Wait a minute!

- Executable scripts run in a subshell

Wait a minute!

- Executable scripts run in a subshell
- They can't change the environment

Wait a minute!

- Executable scripts run in a subshell
- They can't change the environment
- To change the environment needs `source` or `.`, to run in the current shell

Wait a minute!

- Executable scripts run in a subshell
- They can't change the environment
- To change the environment needs `source` or `.,` to run in the current shell
- What even *is* `module`?

Is it an alias?

Slides: git.io/fNEcv

Three thin, dark blue wavy lines that sweep across the bottom of the slide, starting from the left and ending on the right, creating a decorative footer area.

Is it an alias?

```
[s.e.j.bennett@cl2 ~]$ which module
```

Is it an alias?

```
[s.e.j.bennett@cl2 ~]$ which module
```

```
/usr/bin/which: no module in (/usr/lib64/qt-3.3/bin:  
/usr/local/bin:/usr/bin:/usr/local/sbin:/usr/sbin:  
/home/s.e.j.bennett/bin)
```

Is it an alias?

```
[s.e.j.bennett@cl2 ~]$ which module
```

```
/usr/bin/which: no module in (/usr/lib64/qt-3.3/bin:  
/usr/local/bin:/usr/bin:/usr/local/sbin:/usr/sbin:  
/home/s.e.j.bennett/bin)
```

- Nope

Is it an alias?

```
[s.e.j.bennett@cl2 ~]$ which module
```

```
/usr/bin/which: no module in (/usr/lib64/qt-3.3/bin:  
/usr/local/bin:/usr/bin:/usr/local/sbin:/usr/sbin:  
/home/s.e.j.bennett/bin)
```

- Nope
- Where *is* it!?

Aside

```
[s.e.j.bennett@cl2 ~]$ module avail | grep intel
```

Aside

```
[s.e.j.bennett@cl2 ~]$ module avail | grep intel
```

```
----- /apps/modules/legacy -----
```

```
hpcw  raven
```

```
...
```

- module outputs *everything* to stderr

Aside

```
[s.e.j.bennett@cl2 ~]$ module avail | grep intel
```

```
----- /apps/modules/legacy -----
```

```
hpcw  raven
```

```
...
```

- module outputs *everything* to stderr
- Why?

Progress

```
[s.e.j.bennett@cl2 ~]$ cat /etc/profile.d/modules.sh
```

Progress

```
[s.e.j.bennett@cl2 ~]$ cat /etc/profile.d/modules.sh

shell=`/bin/basename `"/bin/ps -p $$ -ocomm=``
if [ -f /usr/share/Modules/init/$shell ]
then
. /usr/share/Modules/init/$shell
else
. /usr/share/Modules/init/sh
fi
```

Eureka

```
[s.e.j.bennett@cl2 ~]$ cat /usr/share/Modules/init/sh
```

Slides: git.io/fNEcv

Three decorative, dark blue wavy lines that sweep across the bottom of the slide, starting from the left and ending on the right.

Eureka

```
[s.e.j.bennett@cl2 ~]$ cat /usr/share/Modules/init/sh  
  
module() { eval ` /usr/bin/modulecmd sh $*`; }  
  
MODULESHOME=/usr/share/Modules  
export MODULESHOME  
  
...
```

- So module is a *function*

Eureka

```
[s.e.j.bennett@c12 ~]$ cat /usr/share/Modules/init/sh  
  
module() { eval ` /usr/bin/modulecmd sh $*`; }  
  
MODULESHOME=/usr/share/Modules  
export MODULESHOME  
  
...
```

- So module is a *function*
- Now, how does it work?

modulecmd

- 7392 lines of Tcl

modulecmd

- 7392 lines of Tcl
- ...

modulecmd

- 7392 lines of Tcl
- ...
- Let's look at a modulefile instead

modulecmd

- 7392 lines of Tcl (<https://git.io/fNEgB>)
- ...
- Let's look at a modulefile instead

Finding some module files

```
[s.e.j.bennett@c12 ~]$ cat /etc/profile.d/02-default-module
```

Finding some module files

```
[s.e.j.bennett@c12 ~]$ cat /etc/profile.d/02-default-modules
```

```
#!/bin/sh
```

```
# Ansible managed
```

```
module use /apps/modules/physics
```

```
module use /apps/modules/medical
```

```
module use /apps/modules/materials
```

```
module use /apps/modules/genomics
```

```
module use /apps/modules/financial
```

```
...
```

Module file directory structure

```
[s.e.j.bennett@c12 ~]$ find /apps/modules/
```

Module file directory structure

```
[s.e.j.bennett@c12 ~]$ find /apps/modules/
```

```
/apps/modules/  
/apps/modules/legacy  
/apps/modules/legacy/raven  
/apps/modules/legacy/hpcw  
/apps/modules/medical  
/apps/modules/materials  
/apps/modules/materials/fluidity  
/apps/modules/materials/OpenFOAM  
/apps/modules/materials/OpenFOAM/5.x-20180613  
/apps/modules/materials/QuantumEspresso  
/apps/modules/materials/QuantumEspresso/6.1  
...
```

OK OK, show me a modulefile already!

```
[s.e.j.bennett@c12 ~]$ more /app/modules/materials/cp2k/2.4
```

OK OK, show me a modulefile already!

```
[s.e.j.bennett@c12 ~]$ more /app/modules/materials/cp2k/2.4
#%Module1.0#####
source /app/modules/system/logger.tcl
##
## CP2K modulefile
##
set appname "cp2k"
set version "2.4"
set para_ver "mpi"
set type "popt"
--More--
```

- Modulefiles are Tcl programs

OK OK, show me a modulefile already!

```
[s.e.j.bennett@c12 ~]$ more /app/modules/materials/cp2k/2.4
#%Module1.0#####
source /app/modules/system/logger.tcl
##
## CP2K modulefile
##
set appname "cp2k"
set version "2.4"
set para_ver "mpi"
set type "popt"
--More--
```

- Modulefiles are Tcl programs
- modulecmd executes these, and does stuff with the result

More of this cp2k modulefile

```
proc ModulesHelp { } {  
    puts stderr "\tLoads PATH and LD_LIBRARY_PATH settings for  
    puts stderr "\tAn example jobscript and input file can be  
    puts stderr "\tThis, along with a README file, is in the  
    puts stderr "\tThe cp2k application directory can be found  
    puts stderr "\tlooking at the output of `which cp2k` and  
    puts stderr "\tMore information on this package can be found  
}  
--More--
```

- Module help gets output to stderr

More of this cp2k modulefile

```
proc ModulesHelp { } {  
    puts stderr "\tLoads PATH and LD_LIBRARY_PATH settings for  
    puts stderr "\tAn example jobscript and input file can be  
    puts stderr "\tThis, along with a README file, is in the  
    puts stderr "\tThe cp2k application directory can be found  
    puts stderr "\tlooking at the output of `which cp2k` and  
    puts stderr "\tMore information on this package can be found  
}  
--More--
```

- Module help gets output to stderr
- Why?

More of this cp2k modulefile

```
proc ModulesHelp { } {  
    puts stderr "\tLoads PATH and LD_LIBRARY_PATH settings for  
    puts stderr "\tAn example jobscript and input file can be  
    puts stderr "\tThis, along with a README file, is in the  
    puts stderr "\tThe cp2k application directory can be found  
    puts stderr "\tlooking at the output of `which cp2k` and  
    puts stderr "\tMore information on this package can be found  
}  
--More--
```

- Module help gets output to stderr
- Why?
- Let's keep going...

Oh no...

```
# Set ulimit unlimited
if { [module-info shelltype sh] } {
    puts "ulimit -s unlimited;"
} elseif { [module-info shelltype csh] } {
    puts "limit stacksize unlimited;"
}
--More--
```

- Why are we putsing shell commands to stdout?

Oh no...

```
# Set ulimit unlimited
if { [module-info shelltype sh] } {
    puts "ulimit -s unlimited;"
} elseif { [module-info shelltype csh] } {
    puts "limit stacksize unlimited;"
}
--More--
```

- Why are we putsing shell commands to stdout?
- Wait

Oh no...

```
# Set ulimit unlimited
if { [module-info shelltype sh] } {
    puts "ulimit -s unlimited;"
} elseif { [module-info shelltype csh] } {
    puts "limit stacksize unlimited;"
}
--More--
```

- Why are we putsing shell commands to stdout?
- Wait
- Surely not

```
module() {  
    eval `/usr/bin/modulecmd sh $*`;  
}
```

`eval` considered harmful (to sanity)

- All actual work is done inside `modulecmd`

`eval` **considered harmful (to sanity)**

- All actual work is done inside `modulecmd`
- The output of which is passed straight to `eval`

`eval` **considered harmful (to sanity)**

- All actual work is done inside `modulecmd`
- The output of which is passed straight to `eval`
- Commands are run by outputting them to `stdout`

`eval` considered harmful (to sanity)

- All actual work is done inside `modulecmd`
- The output of which is passed straight to `eval`
- Commands are run by outputting them to `stdout`
- So messages can only be displayed via `stderr`

`eval` considered harmful (to sanity)

- All actual work is done inside `modulecmd`
- The output of which is passed straight to `eval`
- Commands are run by outputting them to `stdout`
- So messages can only be displayed via `stderr`
- We can cross-check this in the `modulecmd` source...

modulecmd revisited

```
switch -- $::g_shellType {  
  ...  
  {sh} {  
    puts stdout "$var=[charEscaped $::env($var)];\  
    export $var;"  
  }  
  ...  
  {tcl} {  
    set val $::env($var)  
    puts stdout "set ::env($var) {$val};"  
  }  
  {cmd} {  
    set val $::env($var)  
    puts stdout "set $var=$val"  
  }  
  ...  
}
```

Private modules

- Recall module use `/apps/modules/physics`

Private modules

- Recall module use `/apps/modules/physics`
- Last argument is just a path

Private modules

- Recall module use /apps/modules/physics
- Last argument is just a path
- Paths are scanned for modulefile-like things

Private modules

- Recall `module use /apps/modules/physics`
- Last argument is just a path
- Paths are scanned for modulefile-like things
- So we can add a path in `$HOME`

Private modules

- Recall `module use /apps/modules/physics`
- Last argument is just a path
- Paths are scanned for modulefile-like things
- So we can add a path in `$HOME`
- This is built in via the `use.own` module, which looks in `$HOME/privatemodules`

Writing our own module

```
#!/Module -*- tcl -*-  
## This is a module to access something  
proc ModulesHelp { } {  
    puts stderr "This module sets up access to something"  
}  
module-whatis "sets up access to something"  
prereq something/else  
conflict that/other/module  
module load gcc  
setenv      SOMEVERION      0.95  
append-path PATH             /home/[user]/[somedir]/bin  
append-path MANPATH          /home/[user]/[somedir]/man  
append-path LD_LIBRARY_PATH  /home/[user]/[somedir]/lib
```

Thanks for listening!



Slides: git.io/fNEcv